



AMAZON SAGEMAKER PROJECT DOCUMENTATION

A Comprehensive Guide to Training a Logistic Regression Model Using Amazon
Sage Maker

Abstract

This document outlines the steps to train a logistic regression model using Amazon Sage Maker. It covers importing libraries, setting up the Sage Maker session, configuring the estimator, and managing input data. The guide provides clear instructions for setting hyperparameters and starting the training job, offering a concise roadmap for leveraging Sage Maker's capabilities in machine learning.

Romeo Sebola
Romeo.sebola@gmail.com

About the Dataset

The Synthetic Employee Attrition Dataset is designed for analysing and predicting employee turnover. It includes 74,498 samples with detailed employee information such as demographics, job roles, and personal circumstances.

Introduction

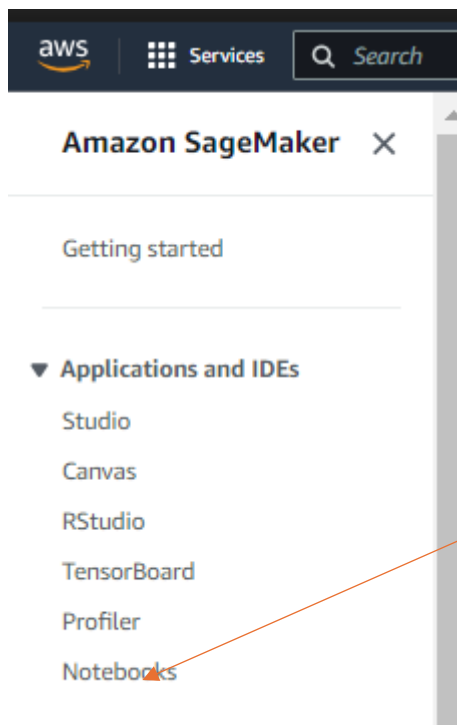
This documentation provides a detailed guide to using Amazon Sage Maker for training a logistic regression model. It includes code snippets and explanations of each step in the process, from data preparation to model training. The goal is to help users understand and execute a machine learning training job using Sage Maker efficiently.

Attrition: Indicates whether the employee has left the company (0 for stayed, 1 for left).

This dataset is valuable for HR analytics and machine learning, helping to understand and predict

Model deployment process

Create Notebook instance



Create notebook instance

Amazon SageMaker provides pre-built fully managed notebook instances that run Jupyter notebooks. The notebook instances include example code for common model training and hosting exercises. [Learn more](#)

Notebook instance settings

Notebook instance name

EmployeeAttrition

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

Notebook instance type

ml.t3.medium

Platform identifier [Learn more](#)

Amazon Linux 2, Jupyter Lab 3

► Additional configuration

Permissions and encryption

IAM role

Notebook instances require permissions to call other services including SageMaker and S3. Choose a role or let us create a role with the [AmazonSageMakerFullAccess](#) IAM policy attached.

AmazonSageMaker-ExecutionRole-20240820T110561

Create role using the role creation wizard

Root access - optional

☒ Enable - Give users root access to the notebook

☐ Disable - Don't give users root access to the notebook

Lifecycle configurations always have root access

Encryption key - optional

Encrypt your notebook data. Choose an existing KMS key or enter a key's ARN.

No Custom Encryption

► Network - optional

Importing Libraries

```
#import the packages or libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pandas.plotting import register_matplotlib_converters
register_matplotlib_converters()
import seaborn as sns
pd.set_option('display.max_columns', None)
import warnings
warnings.filterwarnings('ignore')
```

Loading the Data from the Dataset

```
#Loading the data from the datasets
df=pd.read_csv('train.csv')
df.head(10)
```

| | Employee ID | Age | Gender | Years at Company | Job Role | Monthly Income | Work-Life Balance | Job Satisfaction | Performance Rating | Number of Promotions | Overtime | Distance from Home | Education Level | Marital Status | Number of Dependents |
|---|-------------|-----|--------|------------------|------------|----------------|-------------------|------------------|--------------------|----------------------|----------|--------------------|-------------------|----------------|----------------------|
| 0 | 8410 | 31 | Male | 19 | Education | 5390 | Excellent | Medium | Average | 2 | No | 22 | Associate Degree | Married | 0 |
| 1 | 64756 | 59 | Female | 4 | Media | 5534 | Poor | High | Low | 3 | No | 21 | Master's Degree | Divorced | 3 |
| 2 | 30257 | 24 | Female | 10 | Healthcare | 8159 | Good | High | Low | 0 | No | 11 | Bachelor's Degree | Married | 3 |
| 3 | 65791 | 36 | Female | 7 | Education | 3989 | Good | High | High | 1 | No | 27 | High School | Single | 2 |
| 4 | 65026 | 56 | Male | 41 | Education | 4821 | Fair | Very High | Average | 0 | Yes | 71 | High School | Divorced | 0 |
| 5 | 24368 | 38 | Female | 3 | Technology | 9977 | Fair | High | Below Average | 3 | No | 37 | Bachelor's Degree | Married | 0 |
| 6 | 64970 | 47 | Male | 23 | Education | 3681 | Fair | High | High | 1 | Yes | 75 | High School | Divorced | 3 |
| 7 | 36999 | 48 | Male | 16 | Finance | 11223 | Excellent | Very High | High | 2 | No | 5 | Master's Degree | Married | 4 |

Exploratory Data Analysis (EDA)

```
#EDA
#Exploratory Data Analysis
#view the statistics of the datasets
df.describe()
```

| | Employee ID | Age | Years at Company | Monthly Income | Number of Promotions | Distance from Home | Number of Dependents | Company Tenure |
|-------|--------------|-------------|------------------|----------------|----------------------|--------------------|----------------------|----------------|
| count | 8186.000000 | 8186.000000 | 8186.000000 | 8186.000000 | 8186.000000 | 8186.000000 | 8186.000000 | 8186.000000 |
| mean | 37287.220743 | 38.615075 | 15.667237 | 7336.613609 | 0.827144 | 49.724408 | 1.649890 | 55.863059 |
| std | 21446.665299 | 12.135169 | 11.291536 | 2141.167892 | 0.996414 | 28.518282 | 1.556062 | 25.546607 |
| min | 8.000000 | 18.000000 | 1.000000 | 1855.000000 | 0.000000 | 1.000000 | 0.000000 | 2.000000 |
| 25% | 18757.250000 | 28.000000 | 7.000000 | 5714.000000 | 0.000000 | 25.000000 | 0.000000 | 36.000000 |
| 50% | 37088.500000 | 39.000000 | 13.000000 | 7373.000000 | 0.000000 | 50.000000 | 1.000000 | 56.000000 |
| 75% | 55937.750000 | 49.000000 | 23.000000 | 8879.500000 | 1.000000 | 74.000000 | 3.000000 | 76.000000 |
| max | 74488.000000 | 59.000000 | 51.000000 | 15495.000000 | 4.000000 | 99.000000 | 6.000000 | 127.000000 |

Viewing the First Few Rows of the Dataset

```
df.head()
```

| | Employee ID | Age | Gender | Years at Company | Job Role | Monthly Income | Work-Life Balance | Job Satisfaction | Performance Rating | Number of Promotions | Overtime | Distance from Home | Education Level | Marital Status | Number of Dependents |
|---|-------------|-----|--------|------------------|------------|----------------|-------------------|------------------|--------------------|----------------------|----------|--------------------|-------------------|----------------|----------------------|
| 0 | 8410 | 31 | Male | 19 | Education | 5390 | Excellent | Medium | Average | 2 | No | 22 | Associate Degree | Married | 0 |
| 1 | 64756 | 59 | Female | 4 | Media | 5534 | Poor | High | Low | 3 | No | 21 | Master's Degree | Divorced | 3 |
| 2 | 30257 | 24 | Female | 10 | Healthcare | 8159 | Good | High | Low | 0 | No | 11 | Bachelor's Degree | Married | 3 |
| 3 | 65791 | 36 | Female | 7 | Education | 3989 | Good | High | High | 1 | No | 27 | High School | Single | 2 |
| 4 | 65026 | 56 | Male | 41 | Education | 4821 | Fair | Very High | Average | 0 | Yes | 71 | High School | Divorced | 0 |

Checking for Missing Values in the Dataset

```
: df.isnull().values  
  
: array([[False, False, False, ..., False, False, False],  
        [False, False, False, ..., False, False, False],  
        [False, False, False, ..., False, False, False],  
        ...,  
        [False, False, False, ..., False, False, False],  
        [False, False, False, ..., False, False, False],  
        [False, False, False, ...,  True,  True,  True]])
```

Viewing the Dataset Information

```
#Get the information about the datasets
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8186 entries, 0 to 8185
Data columns (total 24 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Employee ID                          8186 non-null   int64
 1   Age                                  8186 non-null   int64
 2   Gender                              8186 non-null   object
 3   Years at Company                    8186 non-null   int64
 4   Job Role                            8186 non-null   object
 5   Monthly Income                      8186 non-null   int64
 6   Work-Life Balance                   8186 non-null   object
 7   Job Satisfaction                    8186 non-null   object
 8   Performance Rating                  8186 non-null   object
 9   Number of Promotions                8186 non-null   int64
10   Overtime                           8186 non-null   object
11   Distance from Home                  8186 non-null   int64
12   Education Level                     8186 non-null   object
13   Marital Status                      8186 non-null   object
14   Number of Dependents                8186 non-null   int64
15   Job Level                           8186 non-null   object
16   Company Size                        8186 non-null   object
17   Company Tenure                      8186 non-null   int64
18   Remote Work                         8186 non-null   object
19   Leadership Opportunities            8186 non-null   object
20   Innovation Opportunities            8186 non-null   object
21   Company Reputation                  8185 non-null   object
22   Employee Recognition                8185 non-null   object
23   Attrition                          8185 non-null   object
dtypes: int64(8), object(16)
memory usage: 1.5+ MB
```

Analyzing the 'Attrition' Column

```
#Get the count of the number of Employee that stayed or left the company
df['Attrition'].value_counts()
```

```
Attrition
Stayed    4282
Left      3903
Name: count, dtype: int64
```

Analyzing Categorical Columns

```
Gender : ['Male' 'Female']
```

```
Gender
```

```
Male      4500
```

```
Female    3686
```

```
Name: count, dtype: int64
```

```
Job Role : ['Education' 'Media' 'Healthcare' 'Technology' 'Finance']
```

```
Job Role
```

```
Technology  2201
```

```
Healthcare  1852
```

```
Education   1638
```

```
Media        1322
```

```
Finance      1173
```

```
Name: count, dtype: int64
```

```
Work-Life Balance : ['Excellent' 'Poor' 'Good' 'Fair']
```

```
Work-Life Balance
```

```
Good        3089
```

```
Fair         2525
```

```
Excellent   1436
```

```
Poor         1136
```

```
Name: count, dtype: int64
```

```
Job Satisfaction : ['Medium' 'High' 'Very High' 'Low']
```

```
Job Satisfaction
```

```
High         4100
```

```
Very High    1647
```

```
Medium        1612
```

```
Low           827
```

```
Name: count, dtype: int64
```

Encoding Categorical Data


```
#lets convert this attrition to label
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['Attrition']= le.fit_transform(df['Attrition'])
df.head(20)
```


| | Employee ID | Age | Gender | Years at Company | Job Role | Monthly Income | Work-Life Balance | Job Satisfaction | Performance Rating | Number of Promotions | Overtime | Distance from Home | Education Level | Marital Status | Number of Dependents |
|---|-------------|-----|--------|------------------|------------|----------------|-------------------|------------------|--------------------|----------------------|----------|--------------------|-------------------|----------------|----------------------|
| 0 | 8410 | 31 | Male | 19 | Education | 5390 | Excellent | Medium | Average | 2 | No | 22 | Associate Degree | Married | 0 |
| 1 | 64756 | 59 | Female | 4 | Media | 5534 | Poor | High | Low | 3 | No | 21 | Master's Degree | Divorced | 3 |
| 2 | 30257 | 24 | Female | 10 | Healthcare | 8159 | Good | High | Low | 0 | No | 11 | Bachelor's Degree | Married | 3 |
| 3 | 65791 | 36 | Female | 7 | Education | 3989 | Good | High | High | 1 | No | 27 | High School | Single | 2 |
| 4 | 65026 | 56 | Male | 41 | Education | 4821 | Fair | Very High | Average | 0 | Yes | 71 | High School | Divorced | 0 |
| 5 | 24368 | 38 | Female | 3 | Technology | 9977 | Fair | High | Below Average | 3 | No | 37 | Bachelor's Degree | Married | 0 |
| 6 | 64970 | 47 | Male | 23 | Education | 3681 | Fair | High | High | 1 | Yes | 75 | High School | Divorced | 3 |

Creating Boolean Masks for 'Attrition' Column

```
stayed = df.Attrition == 0
left = df.Attrition == 1
```


Plotting Histogram of Ages for Employees Who Stayed

```
plt.hist(df[stayed].Age, alpha=0.5, label='stayed')
plt.title('Age')
plt.xlabel('Age')
plt.ylabel('Samples')
plt.legend()
plt.show()
```



Plotting Histogram of Ages for Employees Who Left

```
plt.hist(df[left].Age, alpha = 0.5, label = 'left')
plt.title('Age')
plt.xlabel('Age')
plt.ylabel('Samples')
plt.legend()
plt.show()
```



Viewing Encoded Class Labels

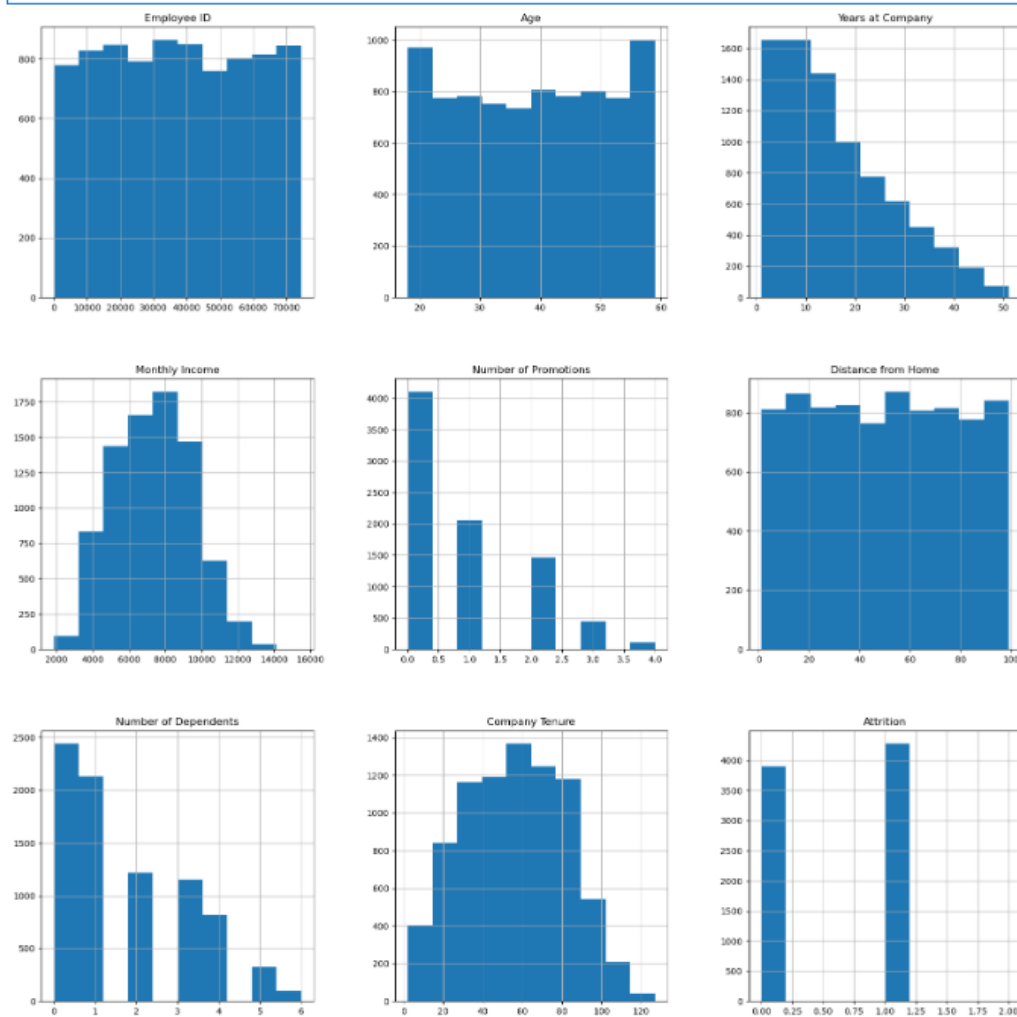
```
: le.classes_
: array(['Left', 'Stayed', nan], dtype=object)
```

Displaying Encoded Class Labels

```
5]: Attrition_Employee = le.classes_
   print(Attrition_Employee)
   ['Left' 'Stayed' nan]
```

Plotting the Distribution of Numerical Columns

```
#plotting the distribution
p = df.hist(figsize = (20,20))
```



Mapping Colors and Plotting Bar Chart for 'Attrition'

```
#check the balance of the data by plotting the count of Attrition
Color_wheel = {1: "#392cf", 2: "#7bc043"}
Colors = df["Attrition"].map(lambda x: Color_wheel.get(x+1))
print(df.Attrition.value_counts())
p = df.Attrition.value_counts().plot(kind="bar")
```

```
Attrition
1    4282
0    3903
2         1
Name: count, dtype: int64
```

Label Encoding Categorical Columns

```
#Define the columns to be label encoded
label_cols = ['Gender', 'Job Role', 'Overtime', 'Education Level', 'Marital Status', 'Company Size', 'Remote Work', 'Leadership Opportunities', 'Innovation Opportunities', 'Work-Life Balance']

#Initialize Label encoders
label_encoders = {col: LabelEncoder() for col in label_cols}

#Apply the Label Encoding
for col in label_cols:
    df[col] = label_encoders[col].fit_transform(df[col])
```

df.head(10)

| | Employee ID | Age | Gender | Years at Company | Job Role | Monthly Income | Work-Life Balance | Job Satisfaction | Performance Rating | Number of Promotions | Overtime | Distance from Home | Education Level | Marital Status | Number of Dependents | Job Level | Company Size | Company Tenure | Remote Work |
|---|-------------|-----|--------|------------------|----------|----------------|-------------------|------------------|--------------------|----------------------|----------|--------------------|-----------------|----------------|----------------------|-----------|--------------|----------------|-------------|
| 0 | 8410 | 31 | 1 | 19 | 0 | 5390 | 0 | 2 | 0 | 2 | 0 | 22 | 0 | 1 | 0 | 1 | 1 | 89 | |
| 1 | 64756 | 59 | 0 | 4 | 3 | 5534 | 3 | 0 | 3 | 3 | 0 | 21 | 3 | 0 | 3 | 1 | 1 | 21 | |
| 2 | 30257 | 24 | 0 | 10 | 2 | 8159 | 2 | 0 | 3 | 0 | 0 | 11 | 1 | 1 | 3 | 1 | 1 | 74 | |
| 3 | 65791 | 36 | 0 | 7 | 0 | 3989 | 2 | 0 | 2 | 1 | 0 | 27 | 2 | 2 | 2 | 1 | 2 | 50 | |
| 4 | 65026 | 56 | 1 | 41 | 0 | 4821 | 1 | 3 | 0 | 0 | 1 | 71 | 2 | 0 | 0 | 2 | 1 | 68 | |
| 5 | 24368 | 38 | 0 | 3 | 4 | 9977 | 1 | 0 | 1 | 3 | 0 | 37 | 1 | 1 | 0 | 1 | 1 | 47 | |
| 6 | 64970 | 47 | 1 | 23 | 0 | 3681 | 1 | 0 | 2 | 1 | 1 | 75 | 2 | 0 | 3 | 0 | 2 | 93 | |
| 7 | 36999 | 48 | 1 | 16 | 1 | 11223 | 0 | 3 | 2 | 2 | 0 | 5 | 3 | 1 | 4 | 0 | 1 | 88 | |
| 8 | 32714 | 57 | 1 | 44 | 0 | 3773 | 2 | 2 | 2 | 1 | 1 | 39 | 2 | 1 | 4 | 0 | 1 | 75 | |
| 9 | 15944 | 24 | 0 | 1 | 2 | 7319 | 3 | 0 | 0 | 1 | 1 | 57 | 4 | 2 | 4 | 0 | 0 | 45 | |

Removing the Last Column and Selecting Remaining Columns

```
x=df.iloc[:, :-1]
#remove the last column Attrition
```

Selecting the Target Column

```
i]: y = df.iloc[:, -1]
```

```
i]: y.head()
```

```
i]: 0    1
     1    1
     2    1
     3    1
     4    1
     Name: Attrition, dtype: int64
```

Splitting Data into Training and Testing Sets

```
from sklearn.model_selection import train_test_split
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X, y, test_size = 0.3)
trainDF = X_Train.join(Y_Train)
```

X_Train.head(5)

| | Employee ID | Age | Gender | Years at Company | Job Role | Monthly Income | Work-Life Balance | Job Satisfaction | Performance Rating | Number of Promotions | Overtime | Distance from Home | Education Level | Marital Status | Number of Dependents | Job Level | Company Size | Company Tenure | Remote Work |
|------|-------------|-----|--------|------------------|----------|----------------|-------------------|------------------|--------------------|----------------------|----------|--------------------|-----------------|----------------|----------------------|-----------|--------------|----------------|-------------|
| 5794 | 15029 | 42 | 1 | 28 | 0 | 4399 | 2 | 0 | 0 | 0 | 0 | 70 | 0 | 1 | 0 | 1 | 2 | 35 | |
| 3116 | 34680 | 39 | 1 | 25 | 2 | 7852 | 1 | 0 | 0 | 1 | 0 | 56 | 4 | 0 | 1 | 1 | 1 | 39 | |
| 1774 | 52289 | 31 | 1 | 4 | 4 | 11030 | 1 | 0 | 0 | 0 | 0 | 25 | 0 | 1 | 2 | 2 | 2 | 25 | |
| 1718 | 22781 | 26 | 0 | 1 | 4 | 10252 | 0 | 0 | 1 | 1 | 0 | 30 | 0 | 0 | 1 | 2 | 0 | 78 | |
| 8135 | 16113 | 47 | 1 | 5 | 2 | 8291 | 1 | 0 | 0 | 2 | 0 | 2 | 0 | 2 | 5 | 2 | 2 | 11 | |

Creating and Inspecting the Training DataFrame

```
trainDF=X_Train.join(Y_Train)
trainDF.head(5)
```

| | Employee ID | Age | Gender | Years at Company | Job Role | Monthly Income | Work-Life Balance | Job Satisfaction | Performance Rating | Number of Promotions | Overtime | Distance from Home | Education Level | Marital Status | Number of Dependents | Job Level | Company Size | Company Tenure | R |
|------|-------------|-----|--------|------------------|----------|----------------|-------------------|------------------|--------------------|----------------------|----------|--------------------|-----------------|----------------|----------------------|-----------|--------------|----------------|---|
| 5794 | 15029 | 42 | 1 | 28 | 0 | 4399 | 2 | 0 | 0 | 0 | 0 | 70 | 0 | 1 | 0 | 1 | 2 | 35 | |
| 3116 | 34680 | 39 | 1 | 25 | 2 | 7852 | 1 | 0 | 0 | 1 | 0 | 56 | 4 | 0 | 1 | 1 | 1 | 39 | |
| 1774 | 52289 | 31 | 1 | 4 | 4 | 11030 | 1 | 0 | 0 | 0 | 0 | 25 | 0 | 1 | 2 | 2 | 2 | 25 | |
| 1718 | 22781 | 26 | 0 | 1 | 4 | 10252 | 0 | 0 | 1 | 1 | 0 | 30 | 0 | 0 | 1 | 2 | 0 | 78 | |
| 8135 | 16113 | 47 | 1 | 5 | 2 | 8291 | 1 | 0 | 0 | 2 | 0 | 2 | 0 | 2 | 5 | 2 | 2 | 11 | |

```
testDF = X_Test.join(Y_Test)
```

Creating and Inspecting the Testing DataFrame

```
testDF = X_Test.join(Y_Test)
testDF.head(5)
```

| | Employee ID | Age | Gender | Years at Company | Job Role | Monthly Income | Work-Life Balance | Job Satisfaction | Performance Rating | Number of Promotions | Overtime | Distance from Home | Education Level | Marital Status | Number of Dependents | Job Level | Company Size | Company Tenure | R |
|------|-------------|-----|--------|------------------|----------|----------------|-------------------|------------------|--------------------|----------------------|----------|--------------------|-----------------|----------------|----------------------|-----------|--------------|----------------|---|
| 3524 | 46223 | 19 | 1 | 5 | 2 | 7423 | 2 | 0 | 0 | 2 | 0 | 94 | 2 | 1 | 2 | 0 | 1 | 48 | |
| 23 | 24208 | 36 | 1 | 13 | 3 | 5874 | 1 | 0 | 0 | 0 | 0 | 16 | 2 | 2 | 1 | 1 | 1 | 40 | |
| 486 | 15082 | 39 | 1 | 7 | 1 | 6991 | 0 | 1 | 0 | 1 | 1 | 65 | 1 | 1 | 1 | 0 | 2 | 84 | |
| 1452 | 72295 | 50 | 0 | 27 | 4 | 8465 | 2 | 0 | 0 | 1 | 1 | 68 | 1 | 2 | 1 | 0 | 1 | 102 | |
| 952 | 4812 | 45 | 0 | 15 | 0 | 4604 | 1 | 0 | 2 | 0 | 0 | 45 | 1 | 2 | 0 | 1 | 2 | 17 | |

Defining a List of Specific Columns

```
column = ['Attrition',
          'Age','Gender',
          'Years at Company',
          'Job Role','Marital Status',
          'Education Level',
          'Job Level',
          'Number of Dependents',
          'Monthly Income',
          'Work-Life Balance',
          'Job Satisfaction','Overtime',
          'Distance from Home','Company Size',
          'Company Tenure','Remote Work',
          'Leadership Opportunities',
          'Innovation Opportunities',
          'Company Reputation',
          'Employee Recognition',
          ]
```

Filtering the Training DataFrame to Specific Columns

```
trainDF = trainDF[column]
trainDF.head(10)
```

| | Attrition | Age | Gender | Years at Company | Job Role | Marital Status | Education Level | Job Level | Number of Dependents | Monthly Income | Work-Life Balance | Job Satisfaction | Overtime | Distance from Home | Company Size | Company Tenure | Remote Work | Leadership Opportunities | Innovation Opportunities |
|------|-----------|-----|--------|------------------|----------|----------------|-----------------|-----------|----------------------|----------------|-------------------|------------------|----------|--------------------|--------------|----------------|-------------|--------------------------|--------------------------|
| 5794 | 1 | 42 | 1 | 28 | 0 | 1 | 0 | 1 | 0 | 4399 | 2 | 0 | 0 | 70 | 2 | 35 | 0 | 0 | 0 |
| 3116 | 1 | 39 | 1 | 25 | 2 | 0 | 4 | 1 | 1 | 7852 | 1 | 0 | 0 | 56 | 1 | 39 | 1 | 1 | 1 |
| 1774 | 1 | 31 | 1 | 4 | 4 | 1 | 0 | 2 | 2 | 11030 | 1 | 0 | 0 | 25 | 2 | 25 | 0 | 0 | 0 |
| 1718 | 0 | 26 | 0 | 1 | 4 | 0 | 0 | 2 | 1 | 10252 | 0 | 0 | 0 | 30 | 0 | 78 | 0 | 0 | 0 |
| 8135 | 0 | 47 | 1 | 5 | 2 | 2 | 0 | 2 | 5 | 8291 | 1 | 0 | 0 | 2 | 2 | 11 | 0 | 0 | 0 |
| 3343 | 1 | 53 | 1 | 32 | 2 | 0 | 1 | 1 | 6 | 6045 | 2 | 0 | 0 | 85 | 2 | 46 | 1 | 0 | 0 |
| 564 | 0 | 36 | 0 | 12 | 4 | 2 | 1 | 1 | 1 | 8656 | 1 | 1 | 1 | 73 | 2 | 71 | 1 | 0 | 0 |
| 773 | 1 | 31 | 1 | 15 | 2 | 0 | 3 | 0 | 0 | 8403 | 2 | 2 | 1 | 71 | 2 | 24 | 0 | 0 | 0 |
| 4213 | 0 | 43 | 1 | 25 | 0 | 2 | 1 | 1 | 0 | 3435 | 0 | 2 | 1 | 7 | 1 | 88 | 0 | 0 | 0 |
| 3354 | 1 | 19 | 1 | 10 | 2 | 1 | 2 | 0 | 2 | 6249 | 2 | 1 | 1 | 85 | 1 | 22 | 1 | 0 | 0 |

25. Filtering the Testing DataFrame to Specific Columns (Excluding the First Column)

```
#Save the trained data#
#write training set#

trainDF.to_csv('traineddataattritions.csv',index=False, index_label='Row',header=False, columns=column)

testDF.head()
```

| | Age | Gender | Years at Company | Job Role | Marital Status | Education Level | Job Level | Number of Dependents | Monthly Income | Work-Life Balance | Job Satisfaction | Overtime | Distance from Home | Company Size | Company Tenure | Remote Work | Leadership Opportunities | Innovation Opportunities |
|------|-----|--------|------------------|----------|----------------|-----------------|-----------|----------------------|----------------|-------------------|------------------|----------|--------------------|--------------|----------------|-------------|--------------------------|--------------------------|
| 3524 | 19 | 1 | 5 | 2 | 1 | 2 | 0 | 2 | 7423 | 2 | 0 | 0 | 94 | 1 | 48 | 0 | 0 | 0 |
| 23 | 36 | 1 | 13 | 3 | 2 | 2 | 1 | 1 | 5874 | 1 | 0 | 0 | 16 | 1 | 40 | 0 | 0 | 0 |
| 486 | 39 | 1 | 7 | 1 | 1 | 1 | 0 | 1 | 6991 | 0 | 1 | 1 | 65 | 2 | 84 | 0 | 0 | 0 |
| 1452 | 50 | 0 | 27 | 4 | 2 | 1 | 0 | 1 | 8465 | 2 | 0 | 1 | 68 | 1 | 102 | 0 | 0 | 0 |
| 952 | 45 | 0 | 15 | 0 | 2 | 1 | 1 | 0 | 4604 | 1 | 0 | 0 | 45 | 2 | 17 | 0 | 0 | 0 |

26. Saving the Filtered Training DataFrame to a CSV File

```
trainDF.to_csv('testddataattritions.csv',index=False, index_label='Row',header=False, columns=column)
```

Importing Libraries for Cloud Integration and Pattern Matching

```
: import boto3 #this package is to integrate with s3 bucket or other cloude service//
import re #this package is to folow a strict pattern to save your work/regular expresession//
```

Specifying Bucket Name and File Paths

```
#Specify bucket name
bucketNM='romeodiabetecbucket'
TrainFile = r'attritiondata/traineddataattritions/traineddataattritions.csv'
TestFile = r'attritiondata/testddataattritions/testddataattritions.csv'
ValFile = r'attritiondata/Val/Val.csv'
ModelFolder = r'attritiondata/model/'
```

Constructing S3 Paths for Data and Model Storage

```
#Loading
s3ModelOutput=r's3://{0}/{1}'.format(bucketNM, ModelFolder)
s3Train=r's3://{0}/{1}'.format(bucketNM, TrainFile)
s3Test=r's3://{0}/{1}'.format(bucketNM, TestFile)
s3Val=r's3://{0}/{1}'.format(bucketNM, ValFile)
```

Constructing the S3 Path for Model Output

```
s3ModelOutput|

's3://romeodiabetecbucket/attritiondata/model/'
```

To document the code for uploading a file to an S3 bucket using boto3, you can structure it like this in your Word document:

Uploading a File to an S3 Bucket

```
: with open('traineddataattritions.csv', 'rb') as f:
    boto3.Session().resource('s3').Bucket(bucketNM).Object(TrainFile).upload_fileobj(f)
```

Uploading a Test Data File to an S3 Bucket

```
with open('testddataattritions.csv', 'rb') as f:
    boto3.Session().resource('s3').Bucket(bucketNM).Object(TestFile).upload_fileobj(f)
```

Importing SageMaker and Getting Execution Role

```
import sagemaker
from sagemaker import get_execution_role
```

Creating a SageMaker Session and Retrieving the Execution Role

```
57]: sagemakerSess=sagemaker.Session()
    role=get_execution_role()
```

Retrieving Docker Image URI for SageMaker Estimator

```
ECRDockercontainer=sagemaker.amazon.amazon_estimator.get_image_uri(sagemakerSess.boto_region_name,'linear-learner','latest')
```

```
WARNING:sagemaker.deprecations:The method get_image_uri has been renamed in sagemaker>=2.
See: https://sagemaker.readthedocs.io/en/stable/v2.html for details.
INFO:sagemaker.image_uris:Same images used for training and inference. Defaulting to image scope: inference.
WARNING:sagemaker.image_uris:Defaulting to the only supported framework/algorithm version: 1. Ignoring framework/algorithm version: latest.
INFO:sagemaker.image_uris:Ignoring unnecessary instance type: None.
```


Configuring a SageMaker Estimator for Logistic Regression Model

```
LogisticModel=sagemaker.estimator.Estimator(image_uri=ECRdockercontainer,
                                             role=role,
                                             train_instance_count=1,
                                             train_instance_type='ml.m4.xlarge',
                                             output_path=s3ModelOutput,
                                             sagemaker_session=sagemakerSess,
                                             base_job_name = 'Logistic-Demo-v1'
                                             )
```

WARNING:sagemaker.deprecations:train_instance_count has been renamed in sagemaker>=2.
See: <https://sagemaker.readthedocs.io/en/stable/v2.html> for details.
WARNING:sagemaker.deprecations:train_instance_type has been renamed in sagemaker>=2.
See: <https://sagemaker.readthedocs.io/en/stable/v2.html> for details.

Setting and Retrieving Hyperparameters for SageMaker Estimator

```
LogisticModel.set_hyperparameters(predictor_type='binary_classifier', mini_batch_size=100)
```

```
LogisticModel.hyperparameters()
```

```
{'predictor_type': 'binary_classifier', 'mini_batch_size': 100}
```

Configuring S3 Input Data for SageMaker Training

```
trainConfig=sagemaker.session.s3_input(s3_data=s3Train,content_type='text/csv')
```

WARNING:sagemaker.deprecations:The class sagemaker.session.s3_input has been renamed in sagemaker>=2.
See: <https://sagemaker.readthedocs.io/en/stable/v2.html> for details.

Starting the Training Job for SageMaker Estimator

```
LogisticModel.fit({'train':trainConfig})
```

INFO:sagemaker:Creating training-job with name: Logistic-Demo-v1-2024-08-20-12-34-10-548
2024-08-20 12:34:10 Starting - Starting the training job...
2024-08-20 12:34:25 Starting - Preparing the instances for training...
2024-08-20 12:34:56 Downloading - Downloading input data...
2024-08-20 12:35:26 Downloading - Downloading the training image.....
2024-08-20 12:36:37 Training - Training image download completed. Training in progress....Docker entrypoint called with argument(s): train
Running default environment configuration script
[08/20/2024 12:37:01 INFO 139985293625152] Reading default configuration from /opt/amazon/lib/python3.8/site-packages/algorithm/resources/default-input.json: {'mini_batch_size': '1000', 'epochs': '15', 'feature_dim': 'auto', 'use_bias': 'true', 'binary_classifier_model_selection_criteria': 'accuracy', 'f_beta': '1.0', 'target_recall': '0.8', 'target_precision': '0.8', 'num_models': 'auto', 'num_calibration_samples': '1000000', 'init_method': 'uniform', 'init_scale': '0.07', 'init_sigma': '0.01', 'init_bias': '0.0', 'optimizer': 'auto', 'loss': 'auto', 'margin': '1.0', 'quantile': '0.5', 'loss_insensitivity': '0.01', 'huber_delta': '1.0', 'num_classes': '1', 'accuracy_top_k': '3', 'wd': 'auto', 'l1': 'auto', 'momentum': 'auto', 'learning_rate': 'auto', 'beta_1': 'auto', 'beta_2': 'auto', 'bias_lr_mult': 'auto', 'bias_wd_mult': 'auto', 'use_lr_scheduler': 'true', 'lr_scheduler_step': 'auto', 'lr_scheduler_factor': 'auto', 'lr_scheduler_minimum_lr': 'auto', 'positive_example_weight_mult': '1.0', 'balance_multiclass_weights': 'false', 'normalize_data': 'true', 'normalize_label': 'auto', 'unbias_data': 'auto', 'unbias_label': 'auto', 'num_point_for_scaler': '10000', 'kvstore': 'auto', 'num_gpus': 'auto', 'num_kv_servers': 'auto', '_log_level': 'info', '_tuning_objective_metric': '', 'early_stopping_patience': '3', 'early_stopping_tolerance': '0.001', '_enable_profiler': 'false'}
[08/20/2024 12:37:01 INFO 139985293625152] Merging with provided configuration from /opt/ml/input/config/hyperparameters.json: {'mini_batch_size': '100', 'predictor_type': 'binary_classifier'}
[08/20/2024 12:37:01 INFO 139985293625152] Final configuration: {'mini_batch_size': '100', 'epochs': '15', 'feature_dim': 'auto', 'use_bias': 'true', 'binary_classifier_model_selection_criteria': 'accuracy', 'f_beta': '1.0', 'target_recall': '0.8', 'target_precision': '0.8', 'num_models': 'auto', 'num_calibration_samples': '1000000', 'init_method': 'uniform', 'init_scale': '0.07', 'init_sigma': '0.01', 'init_bias': '0.0', 'optimizer': 'auto', 'loss': 'auto', 'margin': '1.0', 'quantile': '0.5', 'loss_insensitivity': '0.01', 'huber_delta': '1.0', 'num_classes': '1', 'accuracy_top_k': '3', 'wd': 'auto', 'l1': 'auto', 'momentum': 'auto', 'learning_rate': 'auto', 'beta_1': 'auto', 'beta_2': 'auto', 'bias_lr_mult': 'auto', 'bias_wd_mult': 'auto', 'use_lr_scheduler': 'true', 'lr_scheduler_step': 'auto', 'lr_scheduler_factor': 'auto', 'lr_scheduler_minimum_lr': 'auto', 'positive_example_weight_mult': '1.0', 'balance_multiclass_weights': 'false', 'normalize_data': 'true', 'normalize_label': 'auto', 'unbias_data': 'auto', 'unbias_label': 'auto', 'num_point_for_scaler': '10000', 'kvstore': 'auto', 'num_gpus': 'auto', 'num_kv_servers': 'auto', '_log_level': 'info', '_tuning_objective_metric': '', 'early_stopping_patience': '3', 'early_stopping_tolerance': '0.001', '_enable_profiler': 'false', 'predictor_type': 'binary_classifier'}

Deploying the Trained Model

```
#Deploying the Trained Model
predictModel=LogisticModel.deploy(initial_instance_count=1,instance_type='ml.m4.xlarge',
                                  endpoint_name = 'RomeoEndpoints')
```


INFO:sagemaker:Creating model with name: Logistic-Demo-v1-2024-08-21-07-37-36-202
INFO:sagemaker:Creating endpoint-config with name RomeoEndpoints
INFO:sagemaker:Creating endpoint with name RomeoEndpoints
-----!

Verifying the Deployed Model


RomeoEEndpoints


arn:aws:sagemaker-us-east-1:3399712748200:endpoint/RomeoEEndpoints

8/21/2024, 9:37:37 AM


InService

8/21/2024, 9:42:59 AM

Endpoint configuration



| Name | ARN | Creation time |
|----------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------|-----------------------|
|  RomeoEEndpoints | arn:aws:sagemaker-us-east-1:3399712748200:endpoint-config/RomeoEEndpoints | 8/21/2024, 9:37:37 AM |

Create a lambda function to consume the model using test dataset

romeolambda

Throttle Copy ARN Actions

Function overview info

Diagram Template

+ Add trigger + Add destination

Description
Last modified: 1 hour ago
Function ARN: arn:aws:lambda:us-east-1:339712748200:function:romeolambda
Function URL: info

Code Test Monitor Configuration Aliases Versions

Code source info

Go to Anything (Ctrl-P)

Environment

- romeolambda /
- lambda_function.py

```

1 import boto3
2
3 def lambda_handler(event, context):
4     # TODO implement
5     runtime_client = boto3.client('runtime.sagemaker')
6     endpoint_name = 'EmployeeAttritionEndpoint'
7     sample = '31,1,17,4,2,0,2,1,10310,0,3,1,89,1,60,0,0,0,3,0'
8     response = runtime_client.invoke_endpoint(endpoint_name=endpoint_name,
9                                               ContentType='text/csv',
10                                              Body=sample)
11     result = response['Body'].read().decode('ascii')
12     print(result)
13
14 return {
15     'statusCode': 200,
16     'body': json.dumps('Hello from Lambda!')}
    
```

Upload from

Configure the lambda function

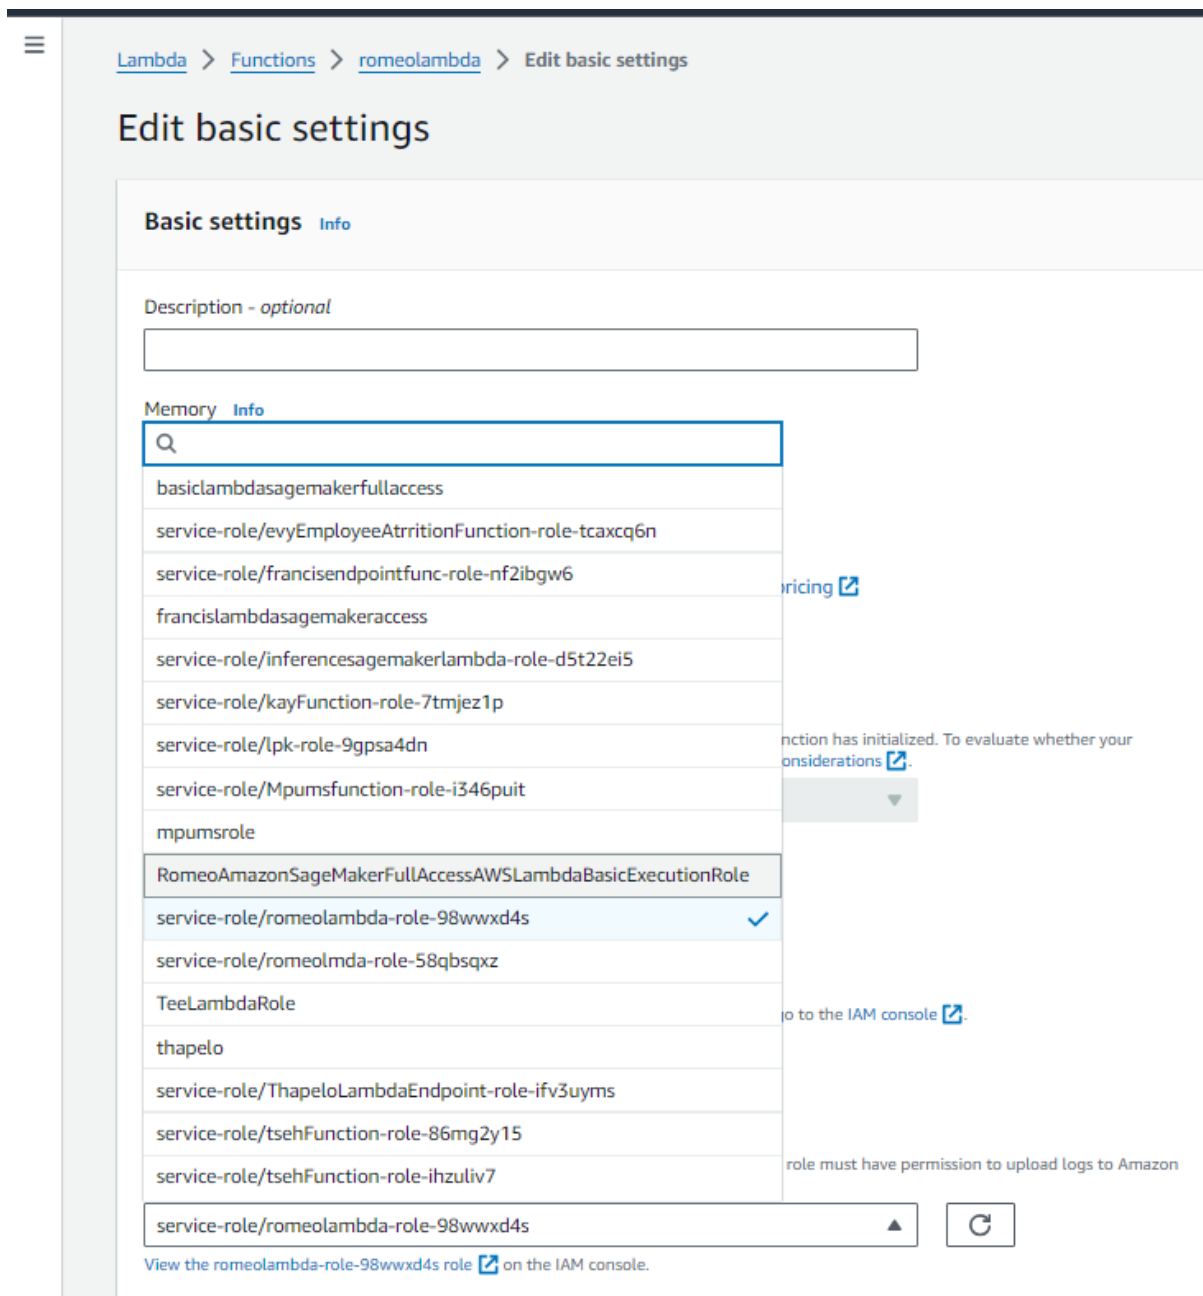

—

Cancel

Save

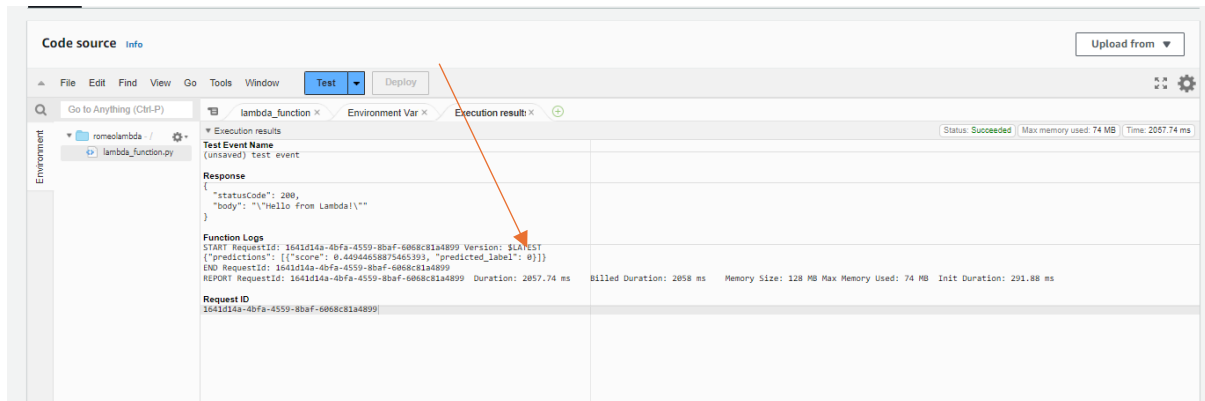
Cancer

Next

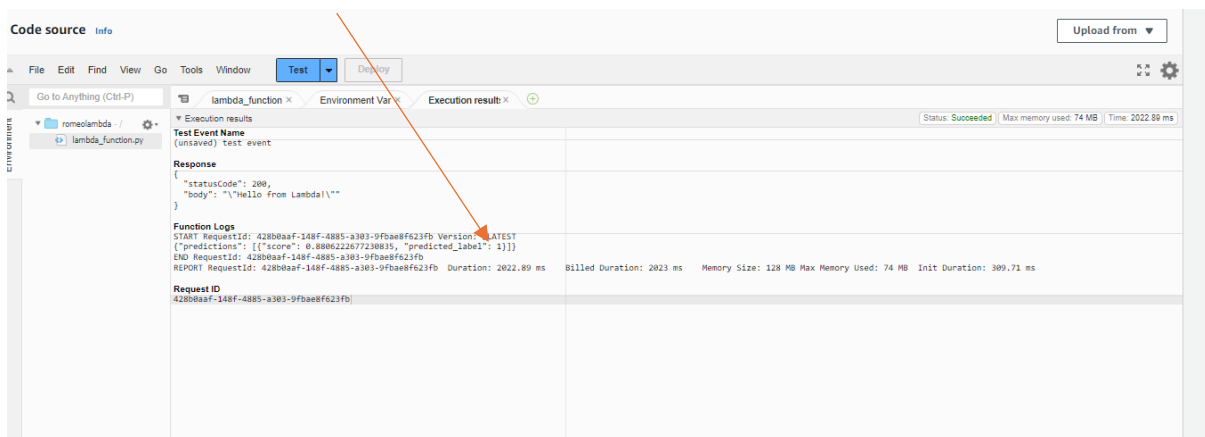


Deploy and Test

-The predicted label indicates “0” That means the Employee will stay



-The predicted label indicates “1” That means the Employee will Leave



Conclusion

In this project, we successfully developed a logistic regression model to predict employee attrition using the Synthetic Employee Attrition Dataset. The process involved multiple steps, from data preparation and exploratory data analysis to model training and deployment using Amazon SageMaker.

This demonstrates the end-to-end process of building, training, and deploying a machine learning model in a cloud environment. The deployed model can now be used for HR analytics to predict employee turnover, providing valuable insights for decision-making. By leveraging Amazon SageMaker, the process was efficient and scalable, showcasing the power of cloud-based machine learning.